

## Practical 1 — Data Loading & Cleaning (Detailed Answers)

### 1. Difference between CSV, Excel, and JSON formats.

CSV: Plain text, comma-separated rows; light, human-readable, universally supported; no data types, no formulas, no multiple sheets.

Excel (.xls/.xlsx): Binary/OOXML; supports multiple sheets, cell types, formulas, styles, filters; heavier; great for business workflows.

JSON: Hierarchical (key-value, lists, nested objects); ideal for APIs and semi-structured data; preserves structure; not columnar by default.

When to use: CSV for tabular exports/ETL; Excel for analyst handoffs; JSON for nested/API data.

### 2. Libraries to load CSV, Excel, JSON into DataFrames

```
pandas.read_csv("file.csv")
pandas.read_excel("file.xlsx", sheet_name="Sheet1")
pandas.read_json("file.json", orient="records")
```

Tip: For big CSVs, add chunksize=100\_000 and iterate; for Excel performance, install engines like openpyxl (xlsx) or xlrd (xls).

### 3. Check missing values after loading

```
df.info()          # quick null counts & dtypes
df.isna().sum()    # per-column null totals
df.isna().mean()*100  # % missing per column
df[df.isna().any(1)].head() # sample rows with any NaN
```

Viva tip: Mention pd.options.display.max\_info\_rows for large sets.

### 4. Handle missing data

Drop: df.dropna() (fast but may lose data).

Impute simple: df.fillna(0) / mean/median/mode per column.

Impute conditional: groupwise: df['col']=df.groupby('grp')['col'].transform(lambda s: s.fillna(s.median()))

Advanced: interpolation (df['col'].interpolate()), KNN/ML imputation (e.g., sklearn.impute.KNNImputer).

Guideline: Pick a method consistent with data generation process; always document imputation.

## 5. Remove duplicates

```
df.duplicated().sum()  
df = df.drop_duplicates()          # all columns  
df = df.drop_duplicates(subset=['id','date'])# key subset
```

Note: If you need first/last occurrence control, use keep='first'/'last'/False.

## 6. Merge multiple datasets

Relational join: pd.merge(a, b, on='key', how='inner|left|right|outer')

Row-wise concat: pd.concat([df1, df2], axis=0)

Column-wise concat: pd.concat([df1, df2], axis=1)

Index join: df1.join(df2, how='left')

Pitfalls: Key dtype mismatches (int vs str), duplicate keys (causing row explosion).

## 7. Purpose of data transformation

To convert raw data into analysis-ready form: type casting, normalization, deriving features, encoding categories, tidying (long↔wide), unit conversions—so models/plots are valid and comparable.

## 8. Compute total sales / average order value (AOV)

```
total_sales = df['Sales'].sum()  
aov = df.groupby('OrderID')['Sales'].sum().mean()  # revenue per order  
# or across time/category:  
daily = df.groupby('Date')['Sales'].sum()  
by_cat = df.groupby('Category')['Sales'].agg(['sum','mean','count'])
```

Tip: Ensure currency/units consistent before aggregating.

## 9. Visualizations for sales trends/product performance

Trends: Line chart (time on x-axis).

Comparisons: Bar/column charts (categories).

Share: Pie/donut (use sparingly).

Distribution: Box/violin (price, order value).

Variability: Error bars; seasonality: line with rolling mean.

Rule: One message per chart; label axes & units.

## 10. Why unify formats before analysis

Consistency reduces parsing errors, simplifies pipelines, ensures accurate joins/aggregations, and speeds EDA/modeling. Example: unify date formats (ISO 8601), currencies, and categorical labels.

## Practical 2 — APIs & JSON (Detailed Answers)

### 1. What is an API & why use it in data analysis?

An API (Application Programming Interface) exposes data/services programmatically (HTTP endpoints). Analysts use APIs to pull fresh, filtered data on demand, enabling automation and reproducibility.

### 2. Purpose of an API key (e.g., OpenWeatherMap)

Identifies the caller, enforces rate limits/quotas, and prevents unauthorized use. Often passed as a query param or header; keep it secret (env vars, not hard-coded).

### 3. Library to send web requests

requests is the go-to: simple, robust, widely used. (Alternatives: httpx async; urllib stdlib.)

### 4. Typical format returned by OpenWeatherMap

JSON—nested objects (e.g., main, weather, wind, coord).

### 5. Extract specific info (temperature/humidity) from JSON

import requests

```
URL = "https://api.openweathermap.org/data/2.5/weather"
params = {"q": "Pune,IN", "appid": API_KEY, "units": "metric"} # metric = °C
res = requests.get(URL, params=params)
data = res.json()
temp_c = data['main']['temp']
humidity = data['main']['humidity']
condition = data['weather'][0]['description']
```

Tip: Use `res.raise_for_status()` and guard keys: `data.get('main', {}).get('temp')`.

### 6. Clean & preprocess weather data from an API

Normalize JSON to columns: `pd.json_normalize(records)`

Units: Kelvin→Celsius ( $^{\circ}\text{C} = \text{K} - 273.15$ ) if not using `units=metric`.

Types: parse timestamps → `pd.to_datetime(dt, unit='s')`; floats for numerics.

Missing: `fillna()` or drop; validate ranges (e.g., humidity 0–100).

Deduplicate: by city+timestamp; timezone alignment.

## 7. Analyze weather patterns

Descriptives: daily/weekly means, max/min, std.

Rolling stats: `df['temp'].rolling(7).mean()` for smoothing.

Anomalies: z-score/IQR on temp or pressure.

Seasonality: group by month or hour-of-day.

## 8. Plots for time-varying weather data

Line: temperature vs time; add rolling mean.

Bar: daily rainfall totals.

Scatter: wind speed vs humidity; temp vs pressure.

Box: monthly temp distribution.

Always label units (°C, mm, m/s, %).

## 9. How geography enhances visualization

Add lat/lon to plot maps (heatmaps/choropleths, bubble maps). Examples: compare city clusters, show temp gradient. With Python, use folium, geopandas, or export to BI tools.

## 10. Advantages of automated API collection vs manual downloads

Real-time/near-real-time updates

Repeatable pipelines (cron/Airflow)

Scalable to many locations/time points

Fewer manual errors; easier auditing/versioning.

## PRACTICAL 3 – Customer Churn Dataset (Detailed Viva Answers)

### 1. What is customer churn? Why is it important to analyze it?

Customer churn means customers stopping the service.

In telecom, churn directly impacts revenue.

Analyzing churn helps companies identify unhappy customers and take steps to retain them (offers, call quality improvements, customer support).

Lower churn = higher customer lifetime value (CLV).

### 2. How to identify missing values in the dataset?

Use pandas functions:

`df.isnull().sum()`

`df.info()`

`df[df.isnull().any(axis=1)]`

Helps locate which columns and how many values are missing.

### 3. Techniques to handle missing data

Delete rows with missing values → `df.dropna()` (but may lose info).

Fill missing values using:

Mean → numeric continuous data.

Median → when outliers exist.

Mode → categorical data.

Group-wise filling: `df.groupby('column').transform(lambda x: x.fillna(x.mean()))`

Interpolation for time-based data.

4. How to detect and remove duplicate records?

`df.duplicated().sum()`

`df = df.drop_duplicates()`

Duplicates cause biased results, especially in churn prediction.

5. Data inconsistencies and how to fix them

Common inconsistencies:

Mixed uppercase/lowercase entries

Extra spaces

Misspelled categorical labels

Fix using:

`df['column'] = df['column'].str.strip().str.lower()`

`df['column'] = df['column'].replace({'yes':'Yes','yess':'Yes'})`

6. Why convert columns to correct data types? How?

Ensures mathematical operations and comparisons behave correctly.

Examples:

`df['TotalCharges'] = df['TotalCharges'].astype(float)`

`df['Churn'] = df['Churn'].astype('category')`

7. What are outliers and how do they affect analysis?

Outliers are values significantly different from the rest.

They can distort averages, produce misleading trends, and reduce model accuracy.

Detect using Boxplot, Z-score, or IQR method.

8. What is feature engineering? Give an example.

Creating new meaningful variables from existing data.

Example in churn dataset:

Create Tenure Group:

`df['TenureGroup'] = pd.cut(df['tenure'], bins=[0,12,24,48,72], labels=['0-1yr','1-2yr','2-4yr','4-6yr'])`

This helps models better capture churn patterns.

9. Why is normalization/scaling important?

To bring all numeric features to similar scale.

Prevents large-scale features from dominating model calculations.

Examples: StandardScaler, MinMaxScaler.

10. How to split data into training and testing sets? Why?

Use `train_test_split()` from `sklearn`.

Training set → learn model, Testing set → evaluate generalization.

Prevents overfitting.

## PRACTICAL 4 – Data Wrangling (Detailed Viva Answers)

1. What is data wrangling and why is it important?

Data wrangling = cleaning, transforming, and organizing raw data.

Ensures data is accurate, structured, and ready for analysis or modeling.

2. How to clean column names in a dataset?

Replace spaces, special characters, and use consistent formatting:

```
df.columns = df.columns.str.strip().str.lower().str.replace(' ','_')
```

Good column names reduce errors in queries and functions.

3. Common strategies to handle missing values

Delete rows/columns with too many NaN.

Fill using:

mean/median (numerical)

mode (categorical)

forward-fill / backward-fill for time-series

Choose method based on data meaning, not convenience.

4. How to merge multiple datasets in pandas?

Use based on structure:

`merge()` → join on key columns

`concat()` → append rows/columns

`join()` → join on index

Correct merging ensures combined dataset has logical consistency.

5. Purpose of filtering and subsetting data

To focus on relevant records.

Example:

```
df[df['price'] > 500000] # Filter
```

```
df[['price','area','location']] # Subset
```

Speeds up analysis and improves insights.

6. Difference between Label Encoding and One-Hot Encoding.

Label Encoding:

Converts categories into numeric values (0,1,2...).

Used when categories have order (e.g., small < medium < large).

One-Hot Encoding:

Creates a new binary column for each category.

Used when categories have no natural order (e.g., red, blue, green).

7. How to calculate average sale price by neighborhood?

```
df.groupby('Neighborhood')['SalePrice'].mean()
```

Helps compare pricing trends across areas.

8. Why handle outliers in housing price data?

Outliers inflate mean price and mislead decision-making.

Removing or capping stabilizes analysis and model accuracy.

9. Functions to detect outliers

Boxplot:

```
df.boxplot(column=['SalePrice'])
```

Z-score: values > 3 or < -3 often considered outliers.

IQR Method: (Q3 + 1.5\*IQR) and (Q1 - 1.5\*IQR) thresholds.

10. Next steps after data wrangling

EDA (Exploratory Data Analysis)

Data visualization

Feature selection

Model training

Evaluation and deployment

## PRACTICAL 5 – Data Visualization using Matplotlib

1. What is Matplotlib and why is it used?

Matplotlib is a Python data visualization library.

It is used to create graphs and charts from data.

Helps convert numerical data into visual patterns for easier understanding.

Useful for EDA (exploratory data analysis) and presentations.

2. How do you import Matplotlib?

```
import matplotlib.pyplot as plt
```

pyplot is used for plotting functions similar to MATLAB style.

plt is a commonly used alias for convenience.

3. Which plot is best to show AQI trends over time?

Line plot is best when the x-axis represents time.

Shows increase/decrease trends clearly.

4. How to plot multiple pollutants (PM2.5, PM10, CO) on the same graph?

```
plt.plot(df['Date'], df['PM2.5'])
```

```
plt.plot(df['Date'], df['PM10'])
```

```
plt.plot(df['Date'], df['CO'])  
plt.show()
```

Plot each pollutant with a separate plt.plot() call.

Add legend to differentiate lines.

#### 5. Purpose of bar plots in AQI analysis

Used to compare pollutant levels across different days, locations, or stations.

Good for category vs quantity comparison.

#### 6. What does a box plot show?

Displays distribution, median, quartiles, and outliers.

Helps detect abnormal AQI spikes.

#### 7. How to add title, axis labels, and legend in Matplotlib?

```
plt.title("AQI Trend")
```

```
plt.xlabel("Date")
```

```
plt.ylabel("AQI Value")
```

```
plt.legend()
```

Title → tells what the graph represents

Labels → clarify axes

Legend → identifies plotted data

#### 8. Use of scatter plots in AQI analysis

Shows relationship between two variables.

Example:

PM2.5 vs AQI → if dots align in a trend, stronger relationship.

#### 9. How to change size and color in Matplotlib?

```
plt.figure(figsize=(10,5))
```

```
plt.plot(df['Date'], df['AQI'], color='red')
```

figsize adjusts plot size.

color changes line color.

#### 10. Why is data visualization important in environmental analysis?

Helps understand pollution trends.

Makes decision-making easier for government & public health bodies.

Identifies peak pollution times and critical zones.

## PRACTICAL 6 – Data Aggregation using Pandas

#### 1. What is data aggregation and why is it important?

Aggregation means summarizing data to get meaningful insights.

Helps convert raw data → useful statistics (sum, mean, max, count).

Useful for business reports and dashboards.

2. Which pandas function is used for aggregation?

groupby() is used to group data based on one or more columns.

Example:

```
df.groupby('Region')['Sales'].sum()
```

3. How to calculate total sales by region?

```
df.groupby('Region')['SalesAmount'].sum()
```

Groups rows by Region and adds up all sales values.

4. Difference between grouping and aggregation (in points)

Grouping:

Splits data into categories (clusters).

Does not perform calculations on its own.

Aggregation:

Applies summary calculations like sum, mean, count on grouped data.

5. How to identify top-performing region?

```
df.groupby('Region')['Sales'].sum().sort_values(ascending=False)
```

Region with the highest value is the top performer.

6. Why group by both region and product category?

To analyze which product sells best in which region.

Helps in strategic marketing and stock distribution.

7. Suitable visualizations for showing sales distribution by region

Bar Chart → compare total sales

Pie Chart → show region-wise contribution to total revenue.

8. What is a stacked bar plot and when is it used?

A bar plot where each bar is divided into segments representing categories.

Used to show total value + category breakdown in one chart.

Example: Sales by Region and Product Type.

9. How to handle missing or inconsistent region names?

Clean text formatting:

```
df['Region'] = df['Region'].str.strip().str.title()
```

Remove leading/trailing spaces.

Standardize name spellings.

10. How does aggregated sales data help business decisions?

Companies can:

Identify high-performing regions

Allocate inventory & marketing budget effectively  
Target low-performing regions for improvement  
Makes decision-making data-driven instead of guesswork.

## PRACTICAL 7 – Time Series Analysis (Stock Market Data)

1. What is time series data and how is it different from other data?

Time series data is data recorded in a sequence over time.

Time (date/hour) plays a major role in analysis.

Values are dependent on previous time values.

Unlike regular datasets, here trend and patterns change with time.

Example: Stock prices recorded daily.

2. Why convert the date column into datetime format in Python?

To enable date-based operations like sorting, filtering, indexing.

Python recognizes datetime format for:

Time-based grouping (day, month, year)

Rolling/moving averages

Time-series plots

`df['Date'] = pd.to_datetime(df['Date'])`

3. Which Python libraries are commonly used for time series analysis?

pandas → loading, cleaning, resampling

matplotlib / seaborn → visualization

statsmodels → forecasting models (ARIMA etc.)

numpy → numeric computations

4. Which plot is used to visualize stock price trends over time?

Line plot

It clearly shows how stock price increases or decreases over days.

`plt.plot(df['Date'], df['Close'])`

5. Purpose of calculating moving averages / rolling averages

To smooth short-term fluctuations.

Highlights long-term trend direction.

`df['MA_7'] = df['Close'].rolling(7).mean()`

Traders use moving averages to identify buy/sell signals.

6. How to detect seasonality in stock price data?

Plot the data over time and observe repeating patterns.

Use seasonal decomposition from statsmodels.

Use autocorrelation plots to check periodic relationships.

7. What does correlation between stock price and volume indicate?

Shows relationship between trading activity and price movement.

If price increases with high volume → strong market interest.

If price rises but volume is low → movement may be weak or temporary.

8. What is the ARIMA model and why is it used?

ARIMA stands for:

Auto Regressive component (AR) → relationship with past values

Integrated (I) → differencing to remove trends

MA (Moving Average) → accounts for past prediction errors

Used for forecasting future values based on past time-series patterns.

9. Key components of time series

Trend → long-term movement (up or down pattern)

Seasonality → repetitive pattern over intervals (daily/weekly/monthly)

Cyclic behavior → irregular movements affected by economy or events

Random noise → unpredictable fluctuations

10. Why is model evaluation important in forecasting? Which metrics are used?

Evaluates how accurate the forecast is.

Helps choose the best performing model.

Common metrics:

MAE (Mean Absolute Error) → average absolute difference

MSE (Mean Square Error) → punishes larger errors

RMSE → square root of MSE for interpretation

MAPE → percentage error, easy to explain in business terms